



DATA SCIENCE BRAIN
@datasciencebrain

DATA

SCIENCE

All Process Cheat Sheet

Save for later reference



DATA LOADING

CSV Files:

```
import pandas as pd
```

Basic CSV read

```
df = pd.read_csv('filename.csv')
```

Specify column names

```
df = pd.read_csv('filename.csv',  
header=None, names=['col1', 'col2'])
```

Specify index column

```
df = pd.read_csv('filename.csv',  
index_col='column_name')
```

Handling Date columns

```
df = pd.read_csv('filename.csv',  
parse_dates=['date_column'])
```

Excel Files:

Basic Excel read

```
df = pd.read_excel('filename.xlsx')
```

Specify sheet name

```
df = pd.read_excel('filename.xlsx',  
sheet_name='Sheet1')
```

Specify column names and index

```
df = pd.read_excel('filename.xlsx', header=0,  
index_col='column_name')
```

Other Formats (JSON, SQL, etc.):

JSON

```
df = pd.read_json('filename.json')
```

SQL Databasefrom sqlalchemy import create_engine

```
engine = create_engine('sqlite:///memory:')  
df = pd.read_sql('SELECT * FROM  
table_name', engine)
```

DATA EXPLORATION & ANALYSIS

General Info:

Display first n rows

```
df.head()
```

Display last n rows

```
df.tail()
```

Data types and non-null counts

```
df.info()
```

Summary statistics

```
df.describe()
```

Descriptive Statistics:

Mean of each column

```
df.mean()
```

Median of each column

```
df.median()
```

Correlation matrix

```
df.corr()
```

Null Values:

Count of null values in each column

```
df.isnull().sum()
```

Drop rows with any null values

```
df.dropna()
```

Fill null values with a specific value

```
df.fillna(value)
```

More Exploratory Data Analysis (EDA):

Value counts for a categorical variable

```
df['column'].value_counts()
```

Unique values in a column

```
df['column'].unique()
```

Cross-tabulation between two columns

```
pd.crosstab(df['column1'], df['column2'])
```

Visualization:

**** LOOK A PART 4 ****



DATA CLEANING

Handling Missing Values:

Interpolate missing values
df.interpolate()

Fill missing values with a specific value
df.fillna(value)

Drop rows with any missing values
df.dropna()

Drop columns with any missing values
df.dropna(axis=1)

Dropping Columns:

Drop columns by name
df.drop(['col1', 'col2'], axis=1, inplace=True)

Drop columns containing a specific pattern
df.drop(df.columns[df.columns.str.contains('pattern')], axis=1, inplace=True)

Data Transformation:

Convert categorical variable to dummy/indicator variables
pd.get_dummies(df['categorical_column'])

Apply a function to each element in a column
df['column'] = df['column'].apply(lambda x: function(x))

Replace values in a column
df['column'].replace({'old_value': 'new_value'}, inplace=True)

Outliers:

Detect and handle outliers (e.g., using Z-score)

from scipy import stats

z_scores = stats.zscore(df['column'])
df_no_outliers = df[(z_scores < 3) & (z_scores > -3)]

DATA VISUALIZATION

Matplotlib for Basic Plots:

import matplotlib.pyplot as plt

Bar chart
plt.bar(df['category'], df['value'], color='blue')

Boxplot
plt.boxplot(df['value'])

Line chart
plt.plot(df['x'], df['y'])

Seaborn for EDA:

import seaborn as sns

Count plot for categorical variable
sns.countplot(x='category', data=df)

Scatter plot matrix
sns.pairplot(df)

Heatmap for correlation
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')

Advanced Plots:

Violin plot
sns.violinplot(x='category', y='value', data=df)

Joint plot for bivariate analysis
sns.jointplot(x='x', y='y', data=df, kind='scatter')

FacetGrid for multi-plot grids
g = sns.FacetGrid(df, col='category', margin_titles=True)
g.map(plt.scatter, 'x', 'y', color='blue')



MACHINE LEARNING

Train-Test Split:

```
from sklearn.model_selection import train_test_split
```

Splitting the data

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Choose a Model:

Linear Models:

- **Linear Regression:** from sklearn.linear_model import LinearRegression
- **Logistic Regression:** from sklearn.linear_model import LogisticRegression

Tree-based Models:

- **Decision Trees:** from sklearn.tree import DecisionTreeClassifier
- **Random Forest:** from sklearn.ensemble import RandomForestClassifier
- **Gradient Boosting:** from sklearn.ensemble import GradientBoostingClassifier

Support Vector Machines:

- **Support Vector Classifier (SVC):** from sklearn.svm import SVC

Nearest Neighbors:

- **K-Nearest Neighbors (KNN):** from sklearn.neighbors import KNeighborsClassifier

Naive Bayes:

- **Gaussian Naive Bayes:** from sklearn.naive_bayes import GaussianNB

Clustering:

- **K-Means:** from sklearn.cluster import KMeans

Neural Networks:

- **Multi-layer Perceptron (MLP):** from sklearn.neural_network import MLPClassifier

Ensemble Methods:

- **AdaBoost:** from sklearn.ensemble import AdaBoostClassifier
- **Bagging:** from sklearn.ensemble import BaggingClassifier

Dimensionality Reduction:

- **Principal Component Analysis (PCA):** from sklearn.decomposition import PCA

Text Processing (for Natural Language Processing):

- **TF-IDF Vectorizer:** from sklearn.feature_extraction.text import TfidfVectorizer
- **Count Vectorizer:** from sklearn.feature_extraction.text import CountVectorizer

MACHINE LEARNING

Train the Model:

Initialize the model

```
model = RandomForestClassifier()
```

Fit the model to the training data

```
model.fit(X_train, y_train)
```

Evaluate the Model:

```
from sklearn.metrics import accuracy_score,
classification_report, confusion_matrix
```

Accuracy score

```
print(accuracy_score(y_test, predictions))
```

Classification report

```
print(classification_report(y_test, predictions))
```

Confusion matrix

```
print(confusion_matrix(y_test, predictions))
```

Regression Evaluation Metrics:

Mean Absolute Error (MAE):

```
from sklearn.metrics import mean_absolute_error
mae = mean_absolute_error(y_true, y_pred)
```

Root Mean Squared Error (RMSE):

```
from sklearn.metrics import mean_squared_error
rmse = mean_squared_error(y_true, y_pred,
squared=False)
```

R-squared (R2 Score):

```
from sklearn.metrics import r2_score
r2 = r2_score(y_true, y_pred)
```

Mean Squared Logarithmic Error (MSLE):

```
from sklearn.metrics import
mean_squared_log_error
msle = mean_squared_log_error(y_true, y_pred)
```

Explained Variance Score:

```
from sklearn.metrics import
explained_variance_score
evs = explained_variance_score(y_true, y_pred)
```

Median Absolute Error:

```
from sklearn.metrics import
median_absolute_error
medae = median_absolute_error(y_true, y_pred)
```



MACHINE LEARNING

Cross-Validation:

```
from sklearn.model_selection import  
cross_val_score
```

Perform cross-validation

```
scores = cross_val_score(model, X, y,  
cv=5)
```

Hyperparameter Tuning:

```
from sklearn.model_selection import  
GridSearchCV
```

Define a parameter grid

```
param_grid = {'n_estimators': [50, 100,  
200], 'max_depth': [None, 10, 20]}
```

GridSearchCV for hyperparameter tuning

```
grid = GridSearchCV(model,  
param_grid, cv=5)  
grid.fit(X_train, y_train)
```

Get the best parameters

```
best_params = grid.best_params_
```

8. Model Saving and Loading:

```
import joblib
```

Save the model

```
joblib.dump(model, 'model.pkl')
```

Load the model

```
loaded_model =  
joblib.load('model.pkl')
```

DEEP LEARNING

Install TensorFlow:

```
pip install tensorflow
```

Import TensorFlow:

```
import tensorflow as tf
```

Build a Neural Network:

```
from tensorflow.keras.models  
import Sequential  
from tensorflow.keras.layers import  
Dense, Dropout, Activation
```

Define a Sequential model

```
model = Sequential()
```

Add layers to the model

```
model.add(Dense(units=128,  
activation='relu', input_shape=  
(input_dim,)))  
model.add(Dropout(0.5))  
model.add(Dense(units=64,  
activation='relu'))  
model.add(Dropout(0.3))  
model.add(Dense(units=output_dim  
, activation='softmax'))
```

Compile the Model:

Compile the model with an optimizer, loss function, and metrics

```
model.compile(optimizer='adam',  
loss='categorical_crossentropy',  
metrics=['accuracy'])
```



DEEP LEARNING

Train the Model:

Train the model with training data

```
model.fit(X_train, y_train, epochs=10,  
batch_size=32, validation_split=0.2)
```

Evaluate the Model:

Evaluate the model on the test data

```
test_loss, test_acc = model.evaluate(X_test,  
y_test)  
print(f'Test Accuracy: {test_acc}')
```

Convolutional Neural Network (CNN):

```
from tensorflow.keras.layers import Conv2D,  
MaxPooling2D, Flatten
```

Example CNN architecture

```
model = Sequential()  
model.add(Conv2D(32, (3, 3), activation='relu',  
input_shape=(img_height, img_width,  
channels)))  
model.add(MaxPooling2D(pool_size=(2, 2)))  
model.add(Conv2D(64, (3, 3), activation='relu'))  
model.add(MaxPooling2D(pool_size=(2, 2)))  
model.add(Flatten())  
model.add(Dense(units=128, activation='relu'))  
model.add(Dense(units=output_dim,  
activation='softmax'))
```

Recurrent Neural Network (RNN):

```
from tensorflow.keras.layers import  
Embedding, LSTM
```

Example RNN architecture for sequence data

```
model = Sequential()  
model.add(Embedding(input_dim=vocab_size,  
output_dim=embedding_dim,  
input_length=max_seq_length))  
model.add(LSTM(units=50,  
return_sequences=True))  
model.add(LSTM(units=50))  
model.add(Dense(units=output_dim,  
activation='softmax'))
```

MODEL SAVING AND LOADING

Model Saving and Loading in scikit-learn:

```
import joblib
```

Save the scikit-learn model

```
joblib.dump(model, 'model.pkl')
```

Load the scikit-learn model

```
loaded_model =  
joblib.load('model.pkl')
```

Model Saving and Loading in TensorFlow:

Save the entire TensorFlow model (including architecture, optimizer, and learned weights)

```
model.save('tensorflow_model')
```

Load a Model:

```
loaded_model =  
tf.keras.models.load_model('tensorflow_model')
```



DEPLOYMENT

scikit-learn Model Deployment:

Once you've trained and saved your scikit-learn model (model.pkl), you can deploy it using various methods depending on your deployment environment.

Flask API:

You can create a simple Flask API to serve your scikit-learn model. Use the flask library to set up an API that receives input data, passes it through the model, and returns predictions.

```
from flask import Flask, request, jsonify
import joblib
```

```
app = Flask(__name__)
model = joblib.load('model.pkl')
```

```
@app.route('/predict', methods=['POST'])def
predict():
    data = request.get_json(force=True)
    prediction = model.predict([data['input']])
    return jsonify({'prediction':
prediction.tolist()})
```

```
if __name__ == '__main__':
    app.run(port=5000)
```

Docker Container:

You can package your Flask API into a Docker container for easy deployment and scalability.

DEPLOYMENT

TensorFlow Model Deployment:

For TensorFlow models, you can use TensorFlow Serving or deploy them as part of a web application.

TensorFlow Serving:

TensorFlow Serving is a system for serving machine learning models in production environments. You can export your TensorFlow model in the SavedModel format and use TensorFlow Serving to deploy it.

Save the TensorFlow model in the SavedModel format

```
model.save('path_to_saved_model',
save_format='tf')
```

Follow the TensorFlow Serving documentation for setting up a server and making predictions.

Flask API for TensorFlow Model:

Similar to scikit-learn, you can use Flask to create an API for serving TensorFlow models.

```
from flask import Flask, request, jsonify
import tensorflow as tf
```

```
app = Flask(__name__)
model =
tf.keras.models.load_model('tensorflow_model')
```

```
@app.route('/predict', methods=['POST'])def
predict():
    data = request.get_json(force=True)
    prediction = model.predict([data['input']])
    return jsonify({'prediction':
prediction.tolist()})
```

```
if __name__ == '__main__':
    app.run(port=5000)
```

